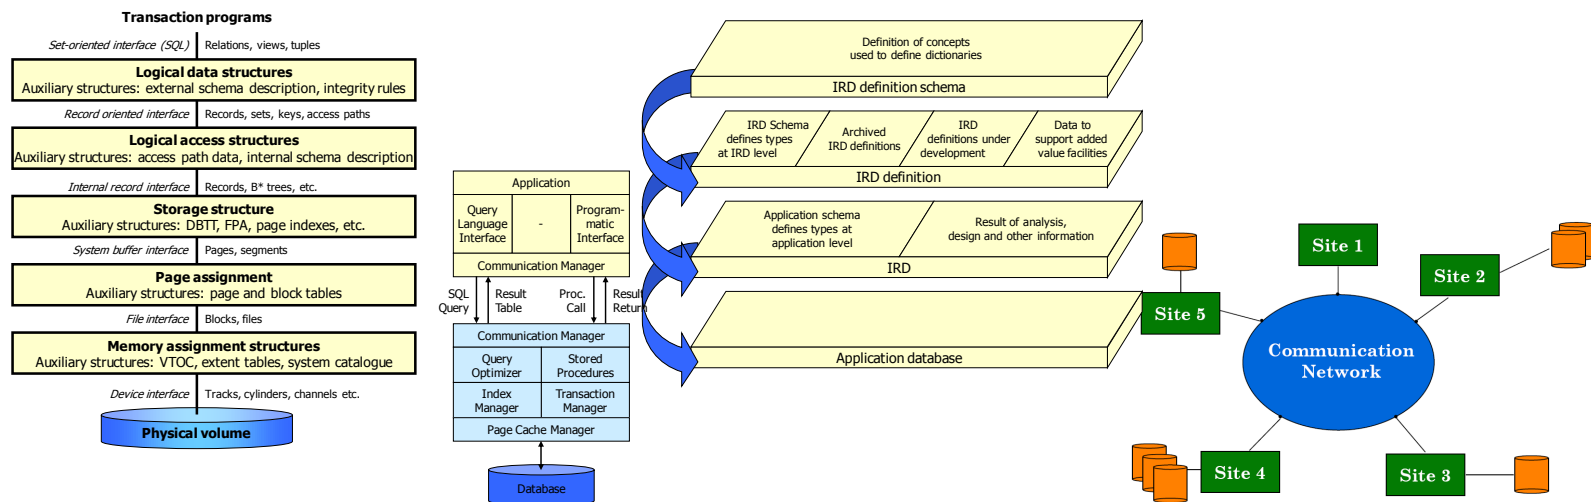


# Chapter 1

## Architecture of Database Systems



# 1. Architecture of Database Systems

Database System = DBMS + Database

- 1.1 Goals and Tasks of DBMS
- 1.2 Basic Architecture of DBMS
- 1.3 DBMS with Transaction Management
- 1.4 Data Dictionary
- 1.5 Distributed Database Systems & Big Data

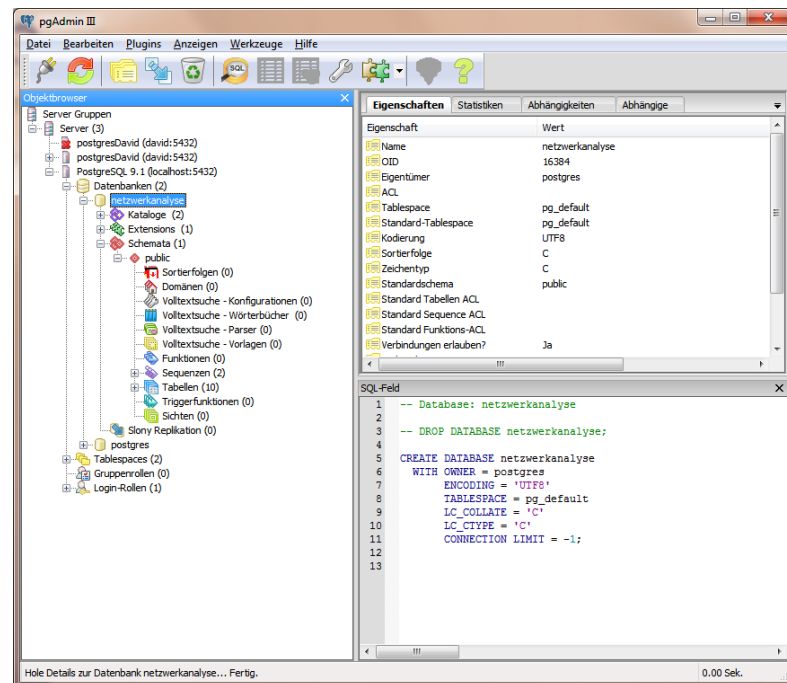
# 1.1 Goals and Tasks of DBMS (1)

- **Data Independence**

- Manage data independent of applications
- Make data available for different applications

- **Data manipulation**

- DML: Data Manipulation Language
  - Retrieve / select
  - Insert / delete / update
- CRUD operations: create – read – update – delete



# 1.1 Goals and Tasks of DBMS (2)

- **Structure definition and integrity assurance**

- DDL: Data Definition Language
- Data dictionary / system catalog / metadata
- Integrity conditions / assertions / constraints

- **Protection of databases in multi-user mode**

- Transaction management
  - **ACID:** Atomicity, Consistency, Isolation, Durability
- Recovery: Restart on error
- Data security and data protection



# 1.1 Goals and Tasks of DBMS (3)

- Performance control
  - Index
  - Clustering / data aggregation
- Realization of user interfaces
  - Interactive end-user interface
  - API: Application Programming Interface



Goals of the first chapter:

- How can a *system architecture* support all the mentioned tasks?
- How can a *system architecture* be designed optimally?

Architecture =  $\langle \{\text{component types}\}, \{\text{connection types}\} \rangle$

	Capacity	Speed	Access time
L1-Cache	16-256 KB	300 GB/s	4 ns
L2-Cache	256 KB-4 MB	300 GB/s	10 ns
L3-Cache	few MBs	~50 GB/s	5-15 ns
RAM	GBs	Up to 10.000 MB/s	10-20 ns
Solid State Disc	GB – TB	200 MB/s	~0.1 ms
Hard disc	TBs	Up to 1.000 MB/s	7 ms
CD/DVD	640 MB – 20 GBs	10 MB/s	150 ms
Streamer	4 GB - >100 GB	2-10 MB/s	100 ms - >10 s
Network	-	1/10/100/1000 MB/s	ca. 1 ms

→ Prefer main memory, reduce disk accesses!

- Databases: Consistent non-volatile memory
- In contrast to OS, DBMS is an application
- Application: Presentation of data, data processing
- Application and DBMS often run on different computers
- Communication: Connections between software systems (is partially a job of OS)
- ISO-OSI Reference Model for communication

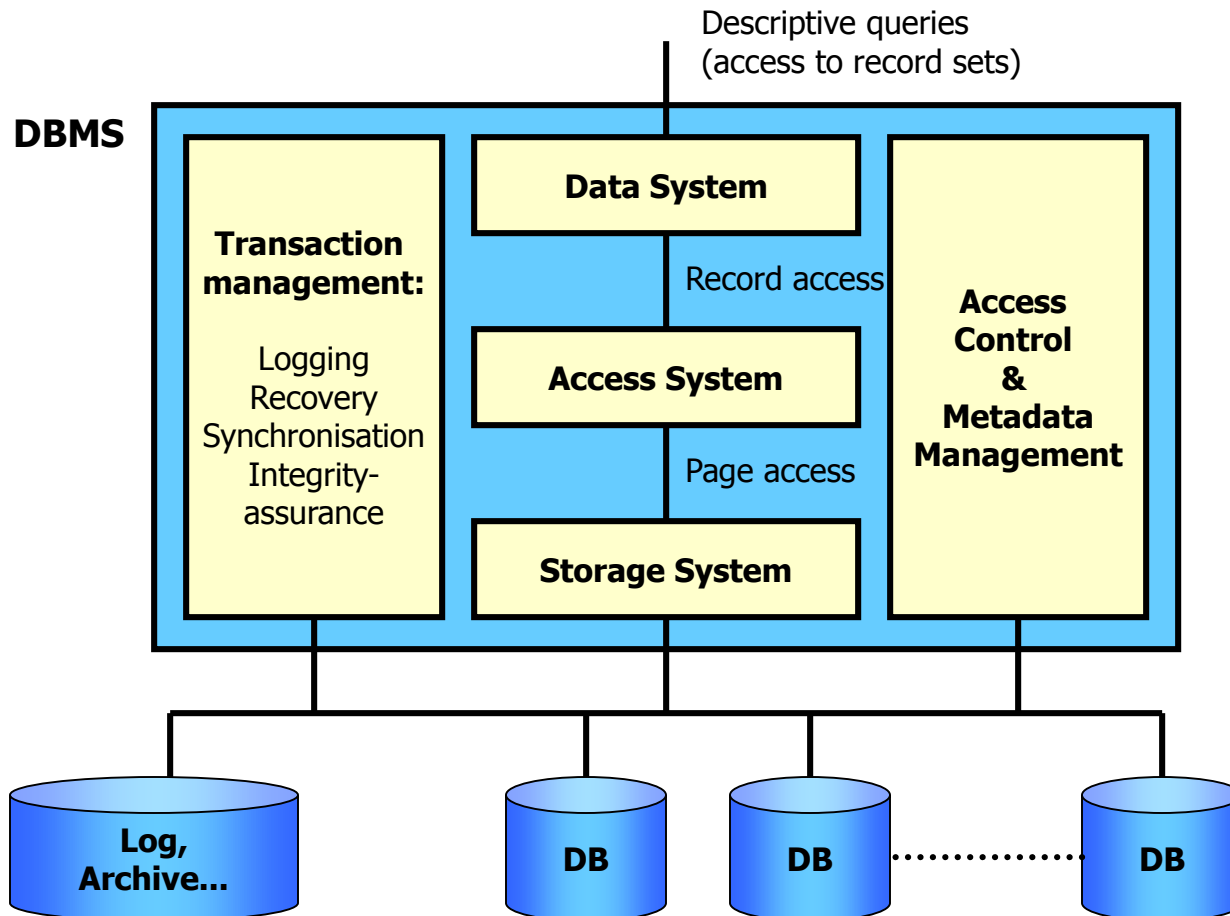
AP	7	Application
	6	Presentation
OS	5	Session
	4	Transport
	3	Network
	2	Data Link
	1	Physical

- Goal of system architecture: **Modularization**
  - Abstraction (concentrate on substantial points)
  - Localization (of procedures & data)
  - Information hiding principle / black box
  - Completeness (on an abstract level)
  - Verifiability
- Concepts for modularization:
  - Functional abstraction
  - Data abstraction
  - Generic modules with objects and methods





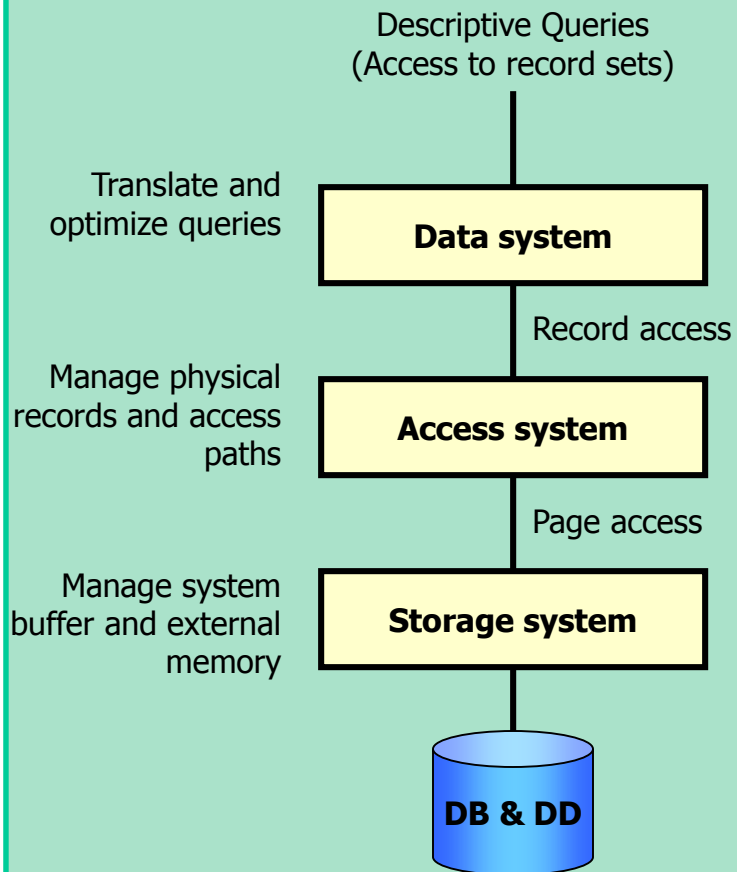
# Simplified Architecture of a DBS



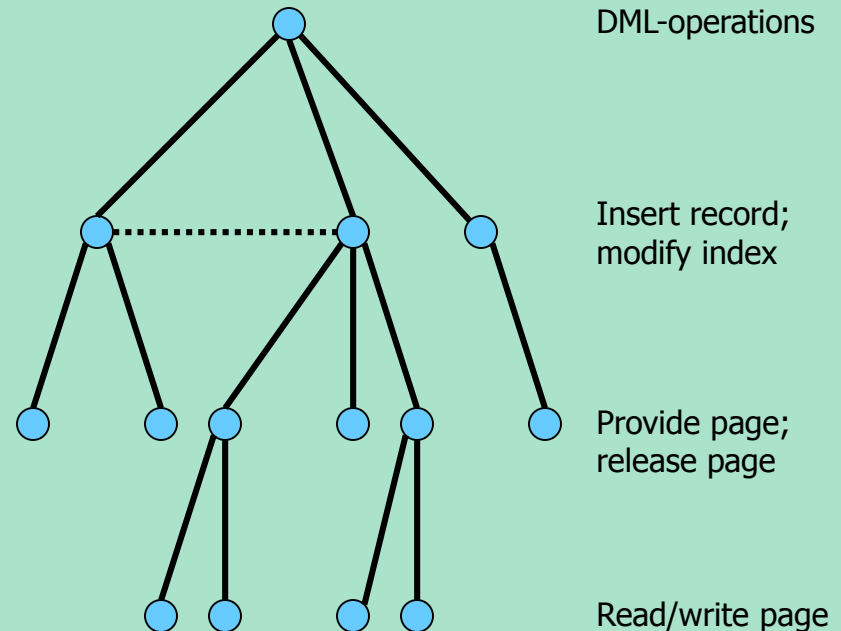
[Härder & Rahm, 2001]

# Components vs. Control Flow

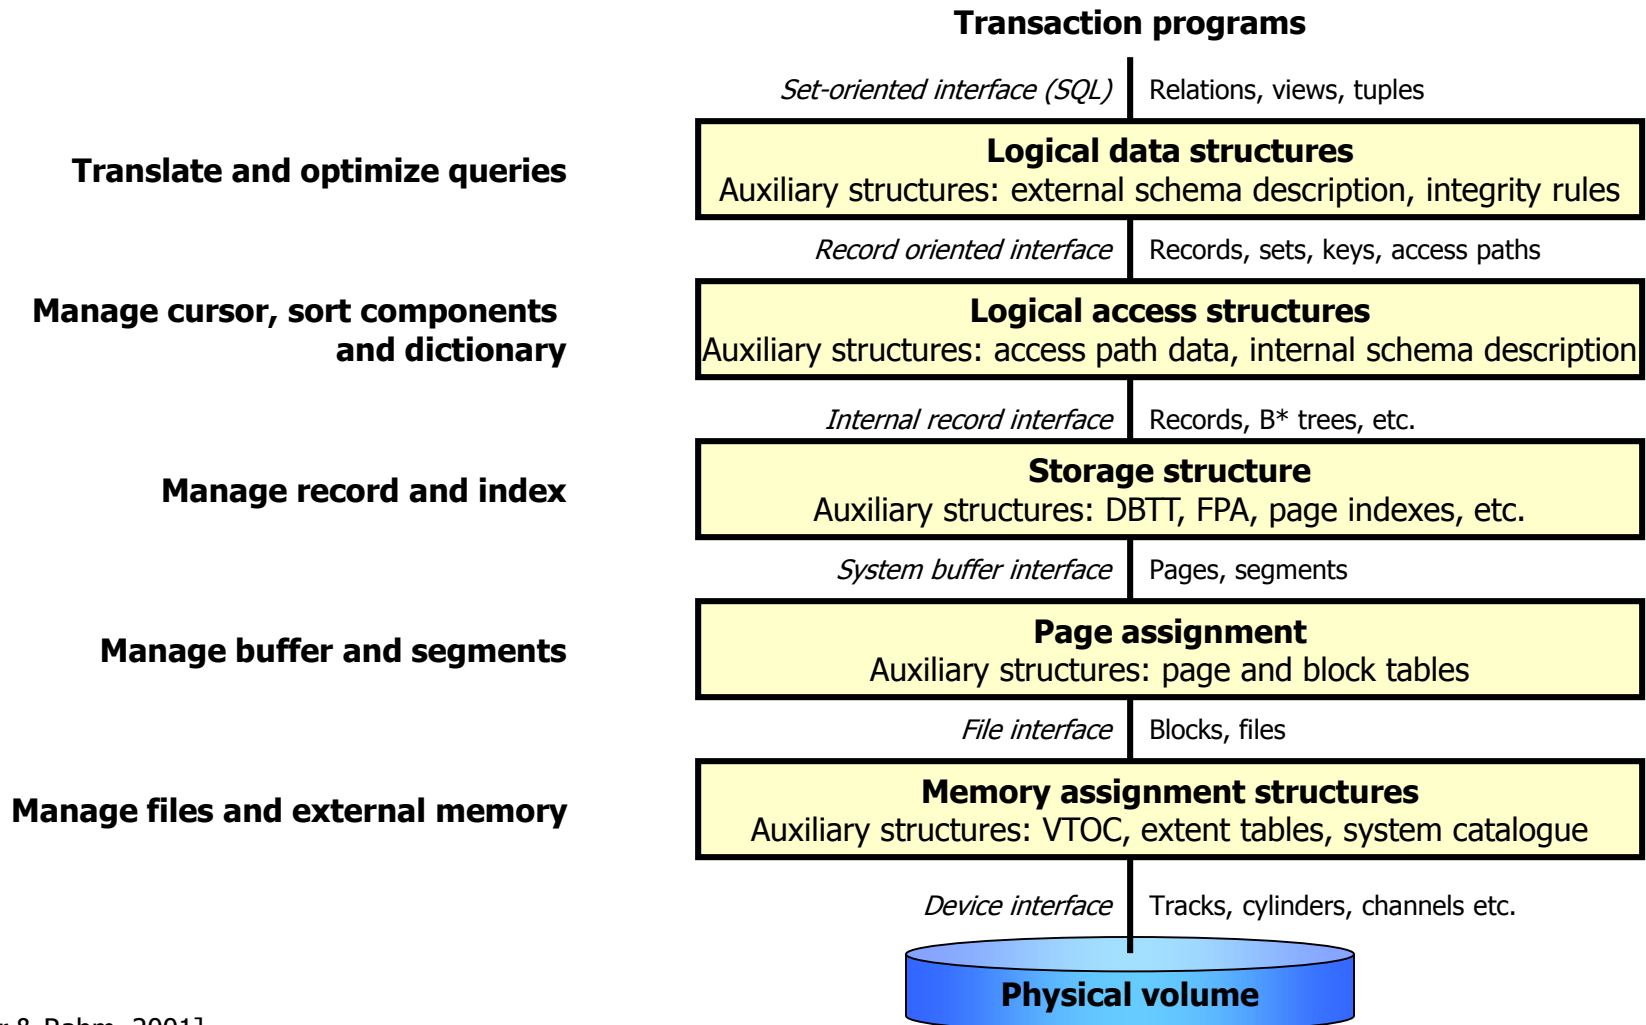
- Components of a DBS



- Dynamic control flow of a DB-operation

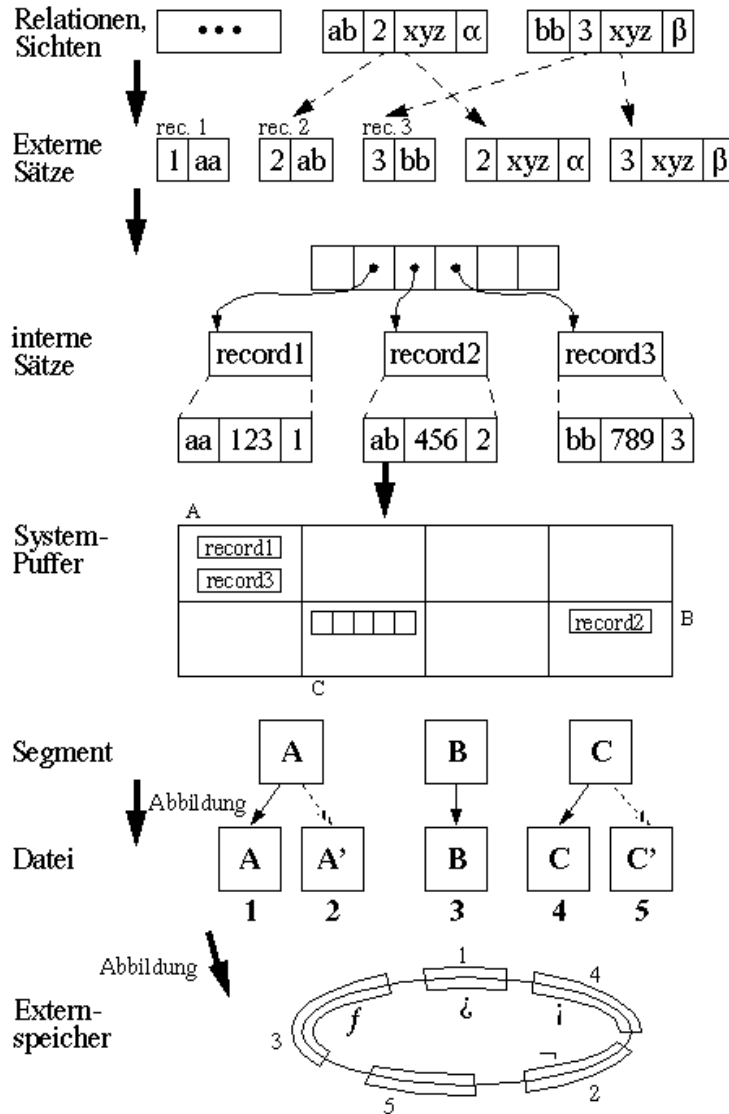


# Five layer model of a DBS



[Härder & Rahm, 2001]

# Example



```
SELECT p.Name FROM Professor p, Vorlesung v
WHERE p.Rang='C4' AND v.Titel='Logik'
AND p.PersNr = v.gelesenVon
```

$$\{ p.name \mid p \in \text{Professor} \wedge \exists v \in \text{Vorlesung} \wedge p.Rang = 'C4' \wedge v.Titel = 'Logik' \wedge p.PersNr = v.gelesenVon \}$$

$$\pi_{Name}(\sigma_{Titel=Logik \wedge Rang=C4}(\text{Professor} \bowtie_{PersNr=gelesenVon} \text{Vorlesung}))$$

$$\pi_{Name}(\sigma_{Rang=C4}(\text{Professor}) \bowtie_{PersNr=gelesenVon} \sigma_{Titel=Logik}(\text{Vorlesung}))$$

```
OPEN CURSOR Vorlesung(Titel='Logik')
FIND NEXT record ...
OPEN CURSOR Professor(Rang='C4')
...
```

```
B+-TREE-SEARCH Vorlesung(Titel='Logik')
FETCH RECORD Vorlesung(...,gelesenVon)
...
B+-TREE-SEARCH Professor(PersNr=gelesenVon)
...
```

```
LOAD PAGE 123
WRITE PAGE 345
...
```

# Data Independence: An Overview


Layer	What is hidden?
Logical data structures	Position indicator and explicit relations in the schema
Logical access paths	Number and kind of the physical access paths; internal representation of records
Storage structures	Management of buffers, logging
Page assignment structures	File mapping, indirect page assignment
Memory assignment structures	Technical features and technical details of external storage media

Problems:

Due to high specialization, functionality of operating system often not usable

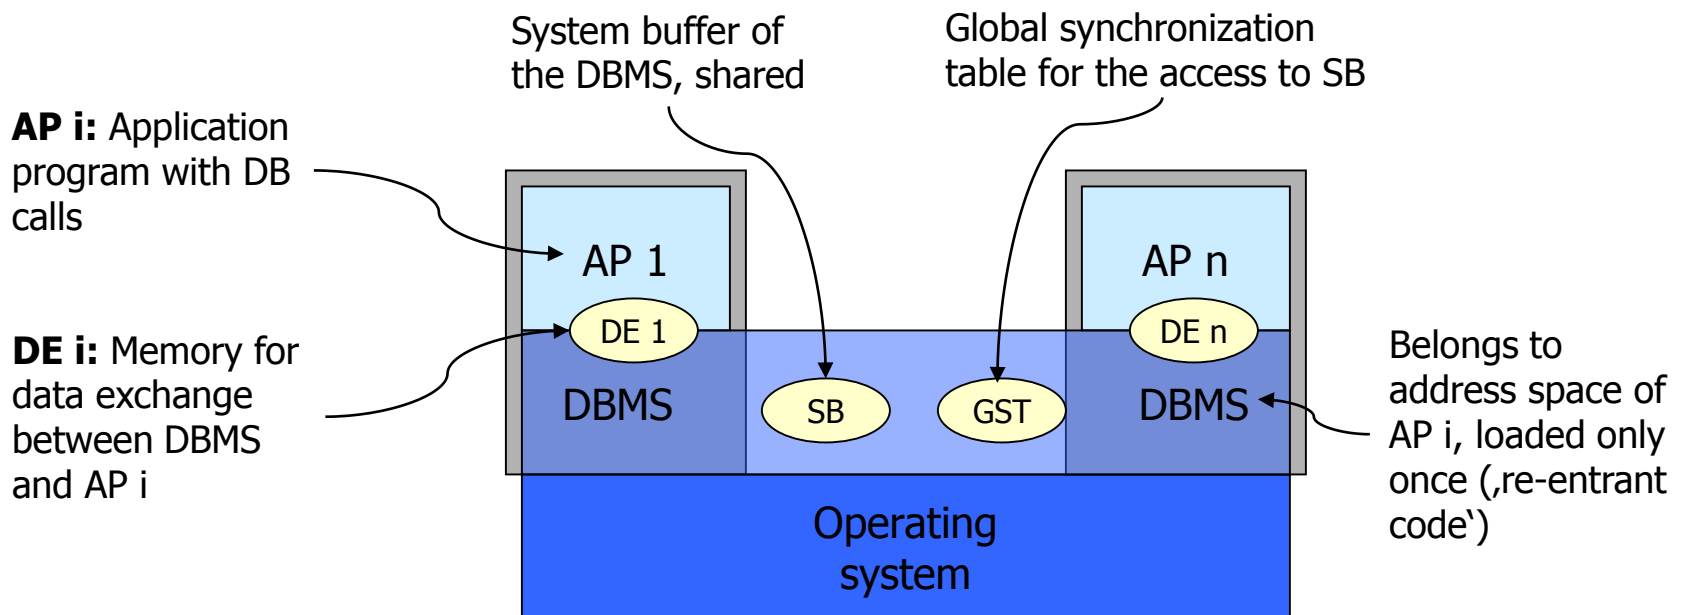
- Segment-file mapping
- Paging
- Shadow memory
- Buffer management
- Dispatching

- Single User DBMS
  - One process, one address space
  - No concurrency control
  - Simple crash recovery
- AP: Application program with DB calls
- DBMS: Belongs to the address space of the AP



Example: PC-database systems (MS Access)


- Multi-user DBMS: Multiple processes, communication via shared address space:
  - Very efficient data exchange via shared memory, but
  - No security concerning errors in the AP



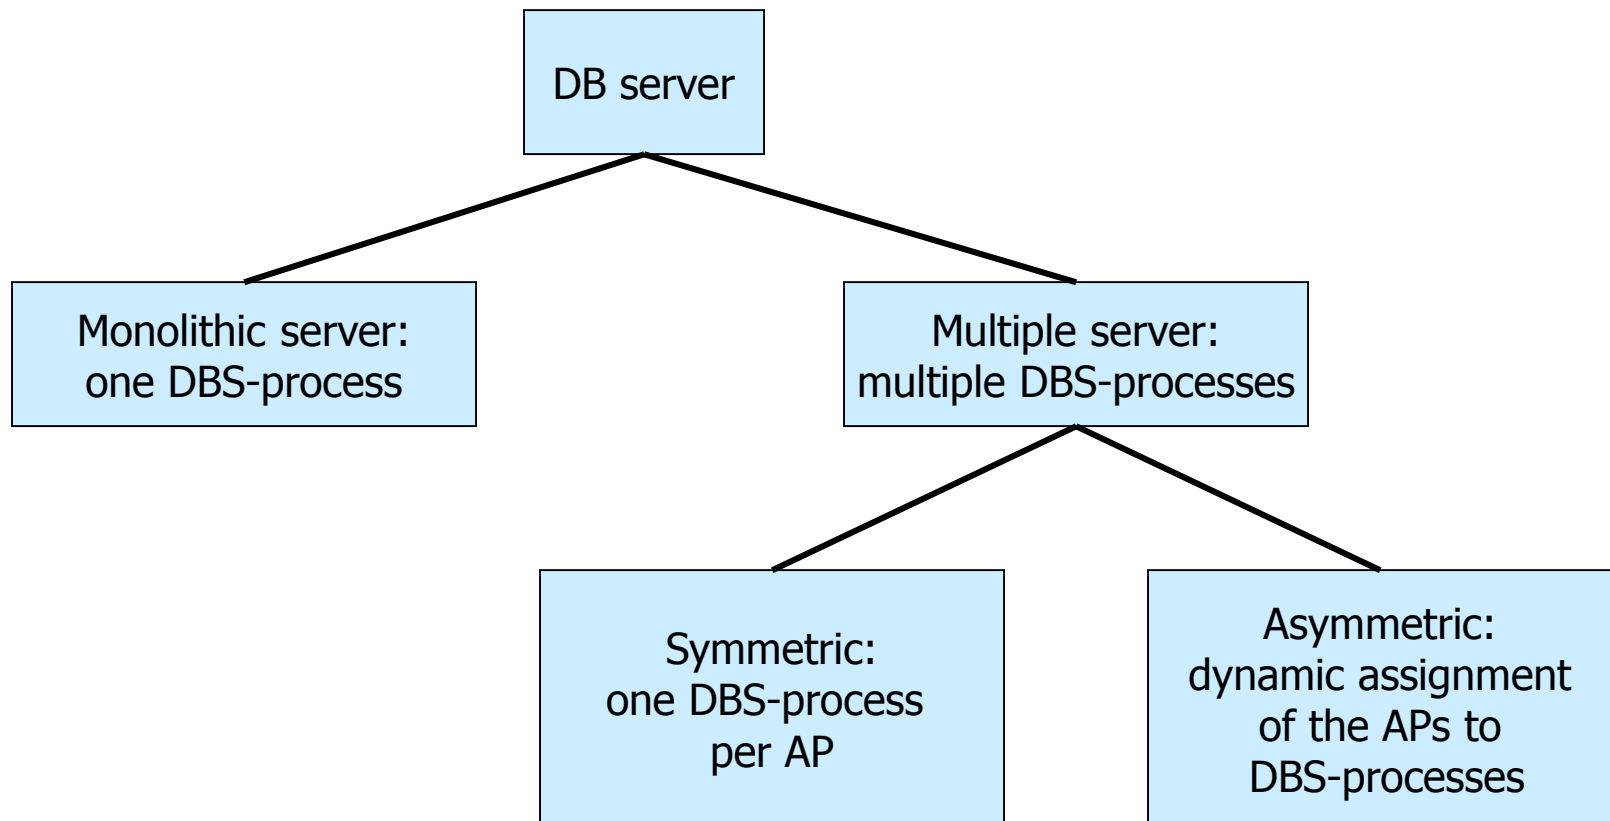
**Example:** IBM System/R\* – research prototype

# Embedded DBMS: 2-Layer-Architecture ("Client/Server")

- Client and server totally separated  $\Rightarrow$  distributed access to DB
- Communication among clients and servers via a network or IPC. Specialized protocols used (JDBC, Net8, TCP/IP)
- Clear separation of client and server







## Multiple DB server processes


- Communication among the servers via “shared memory”
- Communication among clients and servers/dispatcher via OS-mechanisms (IPC) or network software
- **Symmetric assignment:**  
Each client assigned to exact one server process. Static assignment, fixed number  $n$  of servers stated in advance  
⇒ maximal degree of parallelism is  $n$
- **Asymmetric assignment:**  
Each client assigned to a server process by a dispatcher. Fixed number  $n$  of servers stated in advance, but degree of parallelism can be higher.

## Single DB server process

- Synchronized access to system buffer and central system tables
- Server uses multi-threading ("re-entrant code")
- Only one server process for many clients
- DB server process is preferred by OS

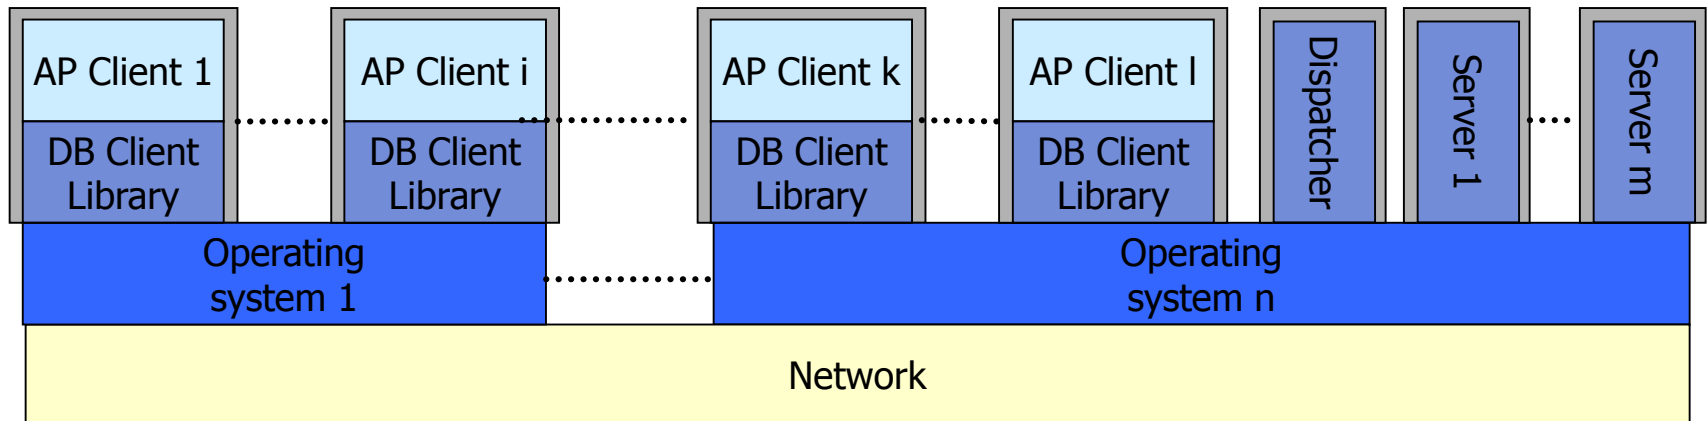
# Monolithic Server

- Own resource management, duplicates OS functions
- Simple communication in the server via shared memory
- Example: PostgreSQL




# Multiple Servers

- DBMS is a compound of different processes
- Communication via operating system or network
- Process scheduling by OS, advantageous in multi-processor computers, because OS manages processor allocation.
- Example: Oracle, IBM DB2, MySQL, SQL Server, Sybase



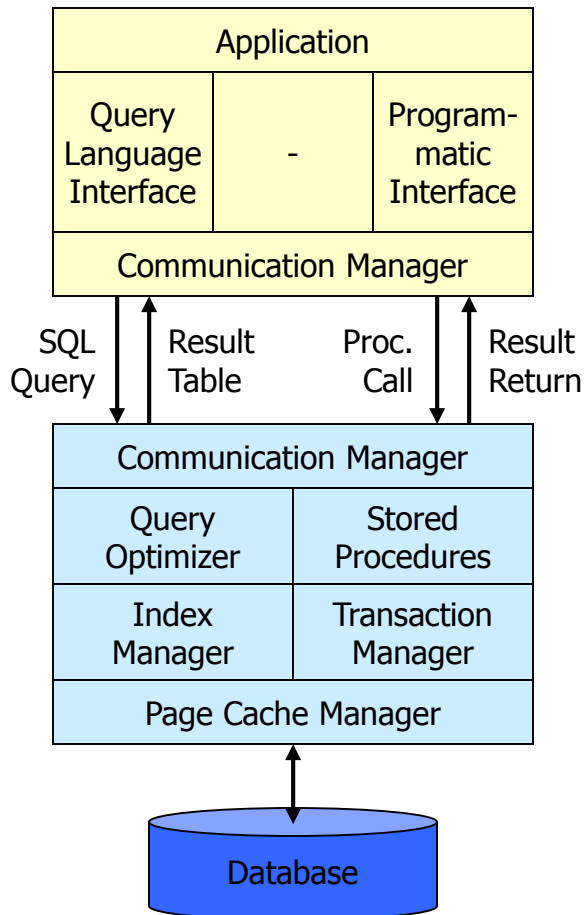
# Tasks of Client / Server

- Tasks of a server:
  - Data management
    - Relation Server
    - Object Server
    - Page Server
  - Application functionality
  
- Tasks of a client:
  - Presentation of data
  - Application functionality

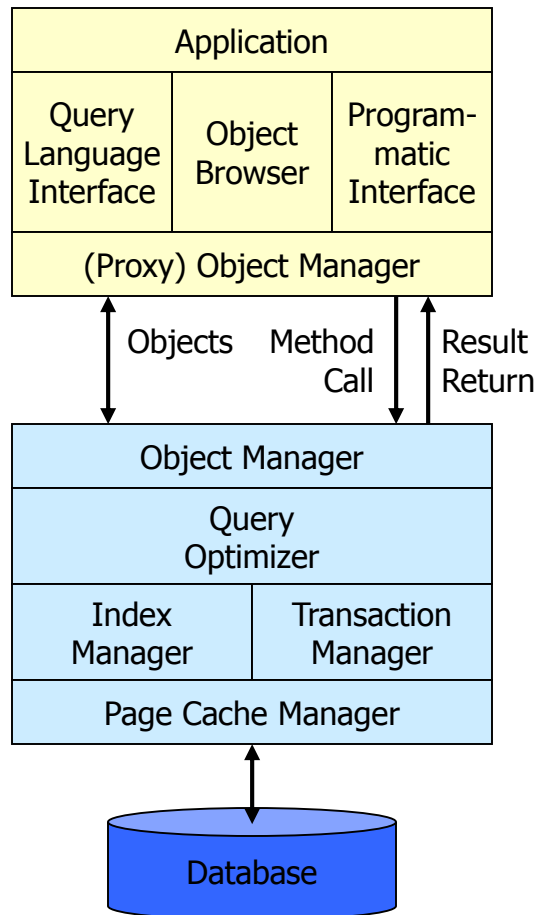


# Server Types

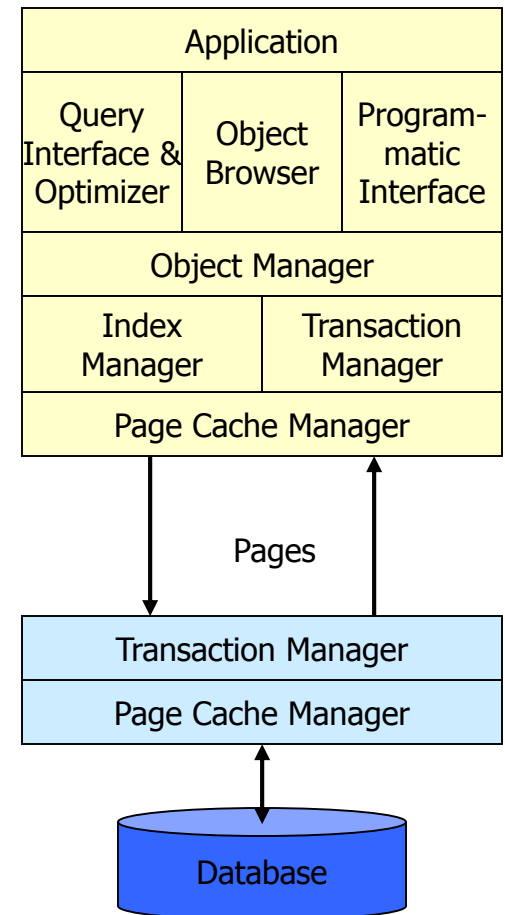
## Relation Server



## Object Server



## Page Server



[Härder & Rahm, 2001]

# 1.3 DBMS with Transaction Management

- Goal:
  - Guarantee correct execution of parallel transactions of multiple users  
 ➔ Details in Chapter 3: Advanced Transaction Management
- Tasks of Transaction Manager (Scheduler):
  - Set locks (e.g., 2 phase locking, 2PL)
  - Detect and solve conflicts (e.g., deadlocks)
- Granularity of Transaction Management (locking level)
  - Synchronisation of access to relations or tuples (relational level)
  - Synchronisation of access to files/blocks (file management level)

*Important: Buffering may cause persistence problems –  
 Conflicts with functionalities of operating system possible!*



# Comparison of Locking Levels


- *Given: Two logically independent transactions T1, T2*
- **Locking on File Level:**  
If T1 and T2 change tuples in the same page, they have to be serialized on the file level.  
⇒ *Inefficient due to unnecessary locks*
- **Locking on Relational Level:**  
T1 and T2 might change tuples in the same block  
➔ Lost Update because T1 and T2 are not isolated  
⇒ *Correctness lost due to logical realization*

*Transaction management can be considered as part of file level (loss of efficiency) or as part of segment layer ( a little bit better) - but:*

Correctness and full performance require comprehensive Transaction management

# 1.4 Data Dictionary

- Problem: Each mapping step means loss of semantics
  - Realisation of operations requires information about
    - Schemas
    - Integrity constraints
    - Index structures
    - Access authorization
    - ...
- Approach: Comprehensive management by data dictionary
  - Internal in DB (uniform model)
  - Stand-alone module (fast and specialized service)



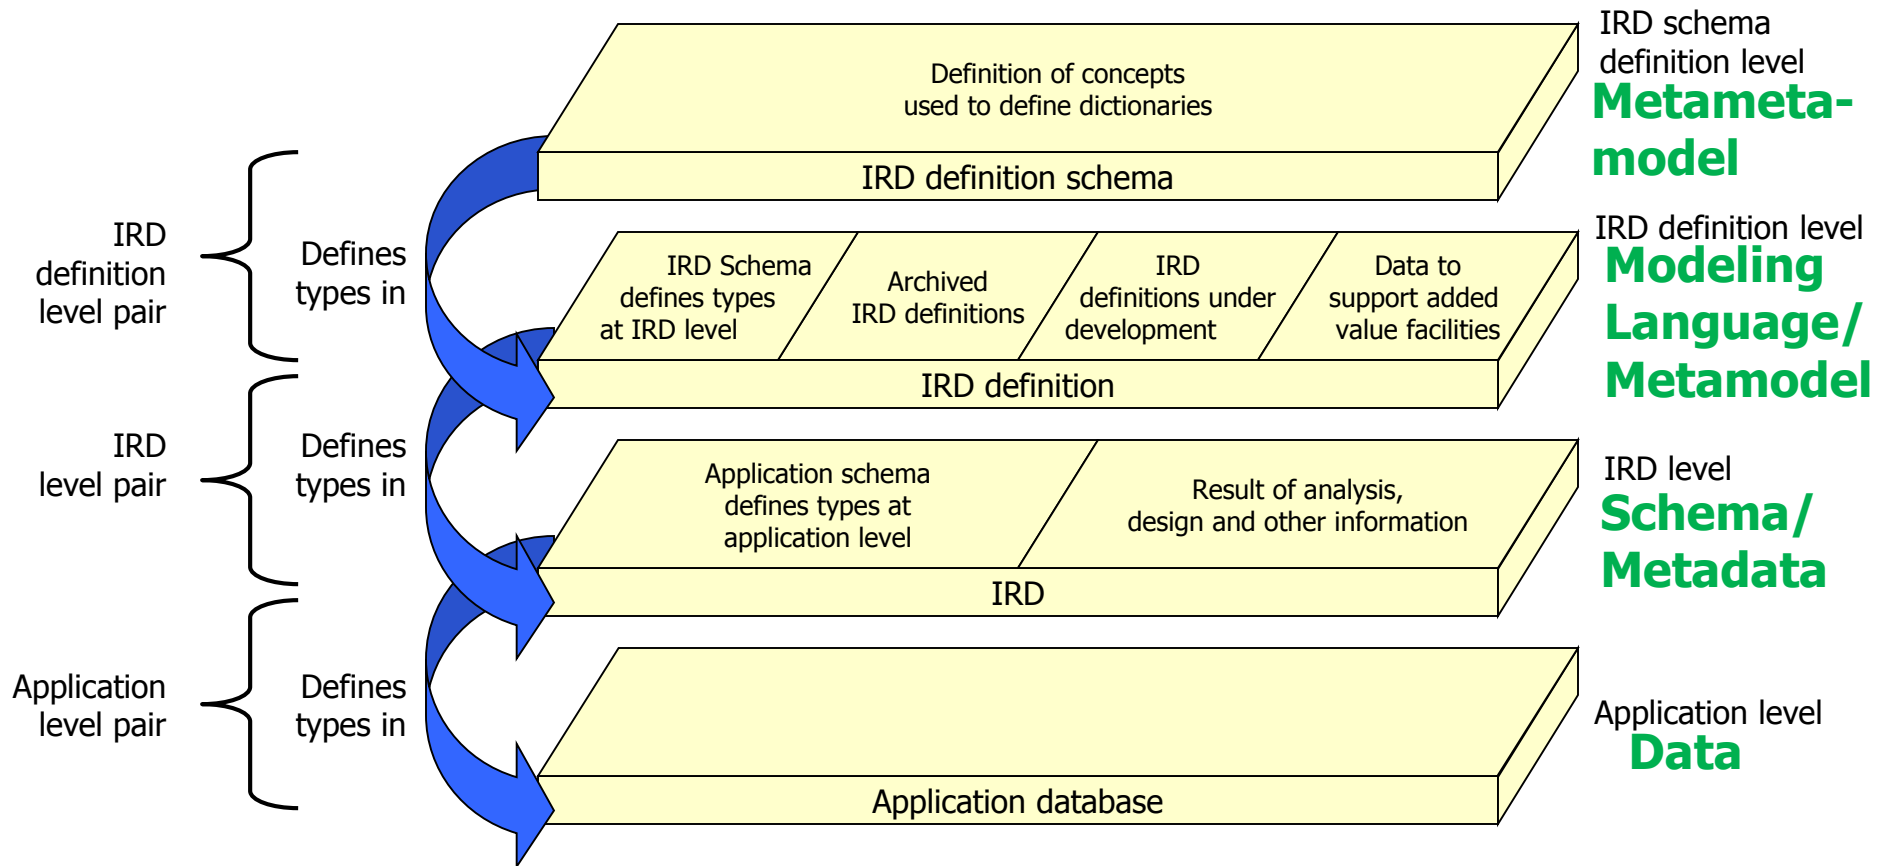
“An Information Resource Dictionary is a [shareable repository](#) for a definition of the information resources relevant to all or part of an enterprise. This may include information about any or all of the following:

- [data](#) needed by the enterprise;
- the computerized and possibly non-computerized [processes](#) which are available for presenting and maintaining such data;
- the available [physical hardware environment](#) on which such data can be represented;
- the [organization](#) of human and physical resources which can make use of the information;
- the human resources responsible for [generating that information](#).

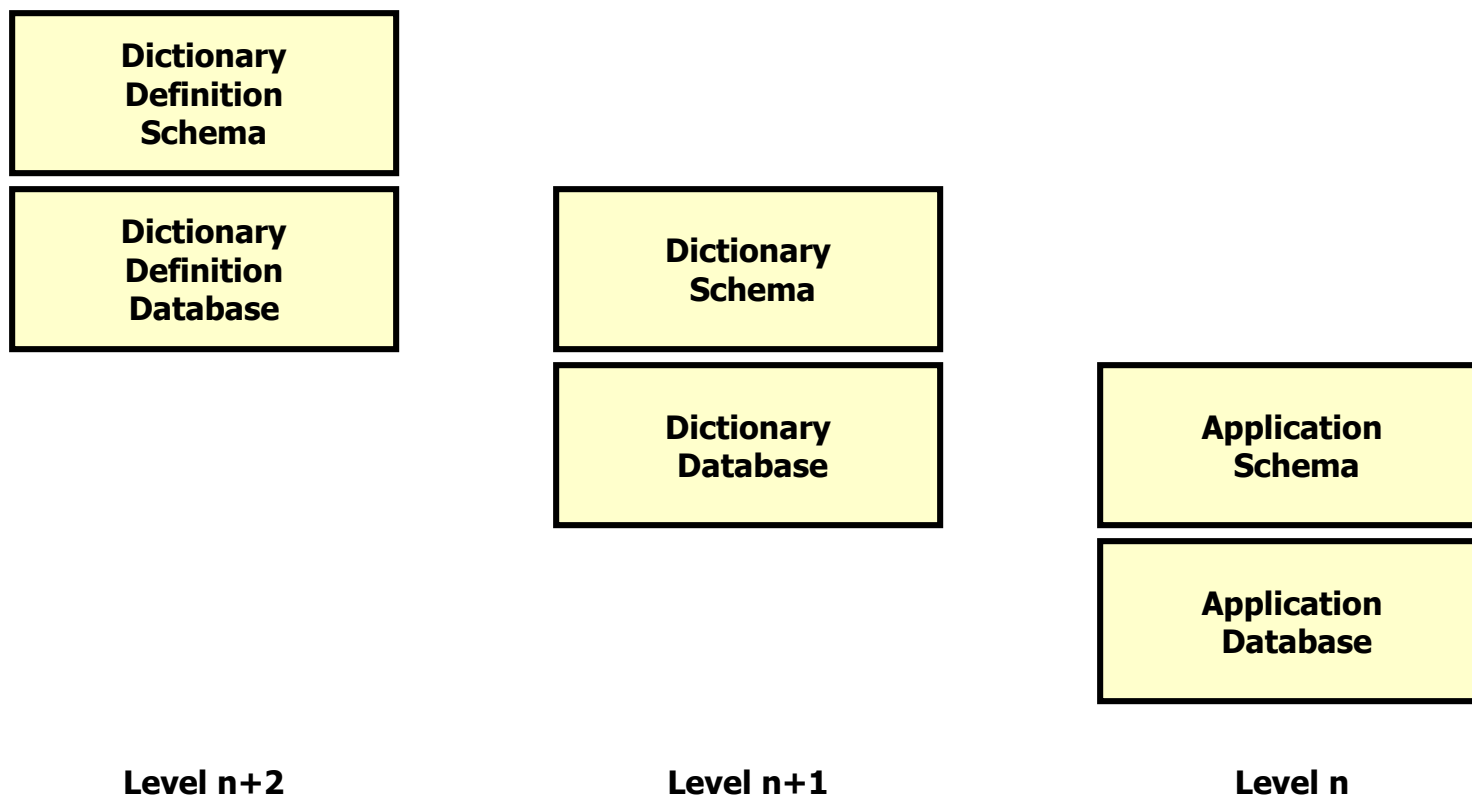
An Information Resource Dictionary System (IRDS) is a system which provides facilities for creating, maintaining and accessing an Information Resource Dictionary (IRD) and its IRD definition.”

[IRDS Framework Standard, ISO 10027:1990]

## IRDS Framework Standard (ISO 10027:1990)




## Example: IRDS Framework with Interlocking Level Pairs



# 1.5 Distributed Database Systems

- A distributed database (DDB) is a collection of multiple, *logically interrelated* databases distributed over a *computer network*
- A distributed database management system (D-DBMS) is the software system that permits the management of the distributed database and makes the distribution *transparent* to the users
- Distributed database system (DDBS)  
DDBS = DDB + D-DBMS



[Özsü & Valduriez, 2011]

# Promises of D-DBMS

- **Transparent** management of distributed, fragmented, and replicated data
- Improved **reliability/availability** through distributed transactions
- Improved **performance**
- Easier and more economical system expansion  
→ **Scalability**




# Implicit Assumptions

- Data stored at a number of sites
  - Each site *logically consists of a single processor.*
- Processors at different sites are interconnected by a computer network
  - no multiprocessors ➡ parallel database systems
- Distributed database is a database, not a collection of files
  - Data is logically related as exhibited in the users' access patterns  
➡ relational data model
- D-DBMS is a full-fledged DBMS
  - ➡ not remote file system, not a Transaction Processing system




# Transparency

- Transparency is the separation of higher level semantics of a system from the lower level implementation issues
- Fundamental issue is to provide **data independence** in the distributed environment
  - ➡ Network (distribution) transparency
  - ➡ Replication transparency
  - ➡ Fragmentation transparency
    - horizontal fragmentation: selection
    - vertical fragmentation: projection
    - hybrid



# D-DBMS: Implementation Alternatives



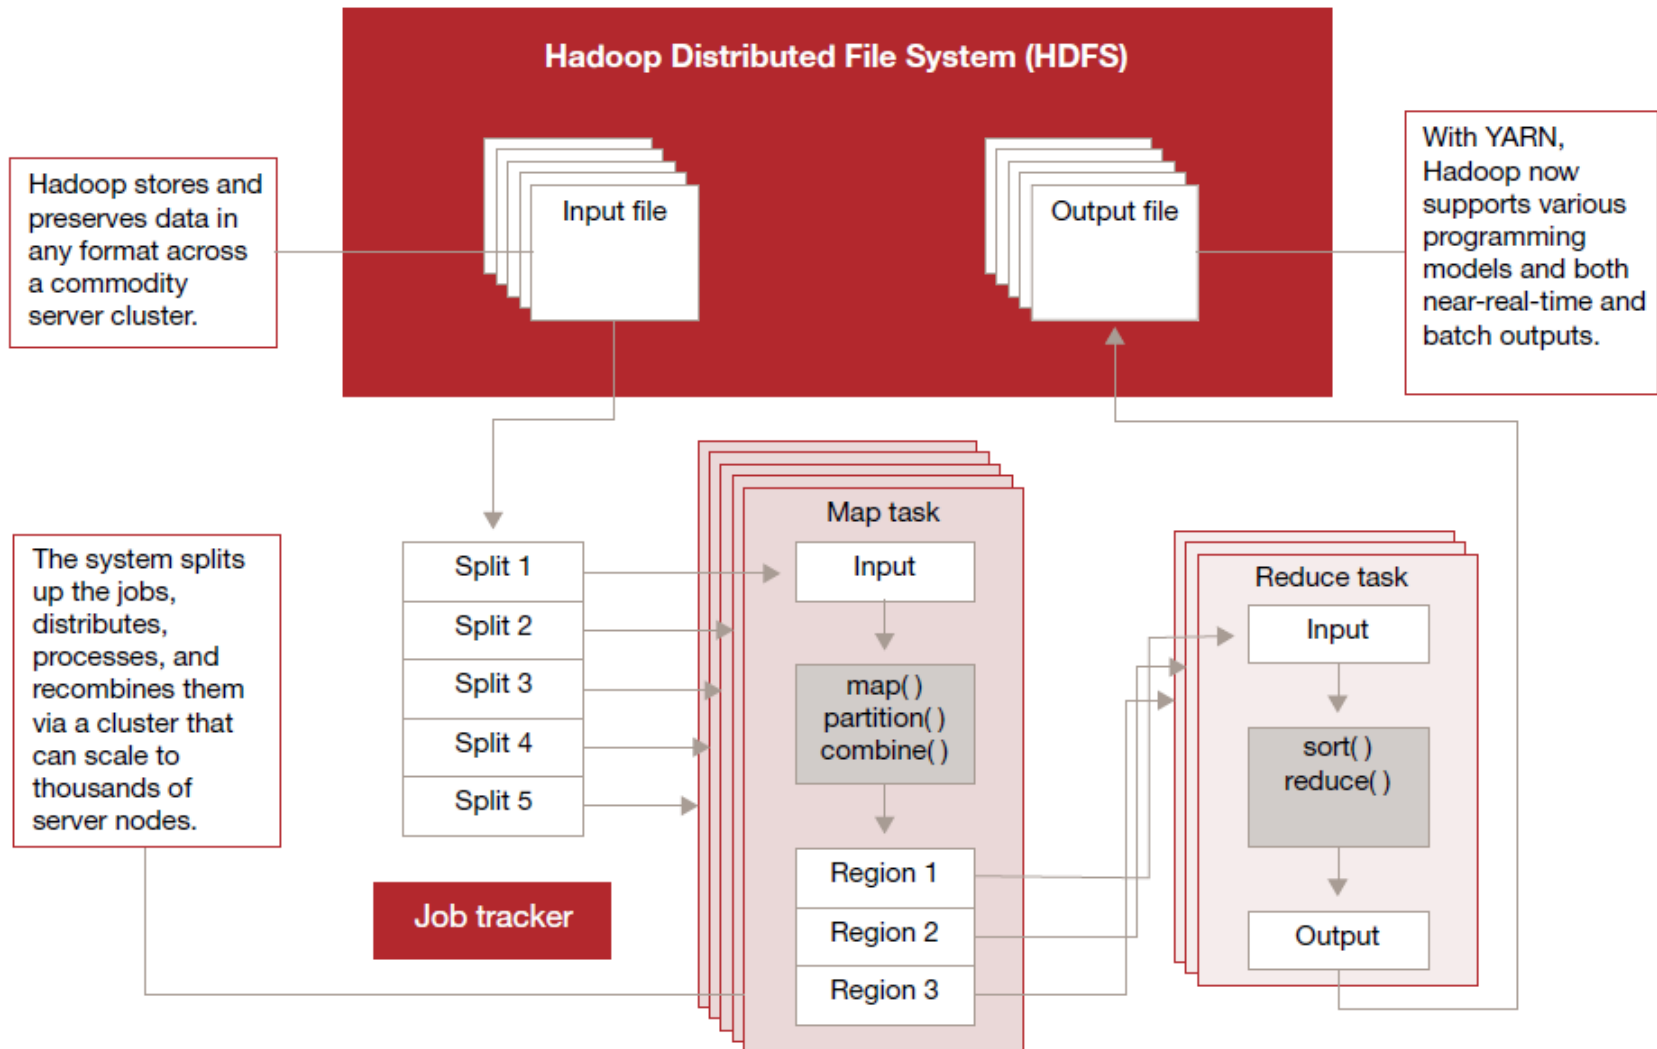
# Big Data Architectures

- Big Data requires large, scalable, distributed architectures
- Four Vs: Volume, Velocity, Variety, Veracity
- Heterogeneity
  - Sources
  - Systems
  - Requirements
  - Client Applications

➔ Big Data Systems are complex eco-systems in which several independent components have to be integrated

➔ Hadoop is not a Big Data system, it is just a component

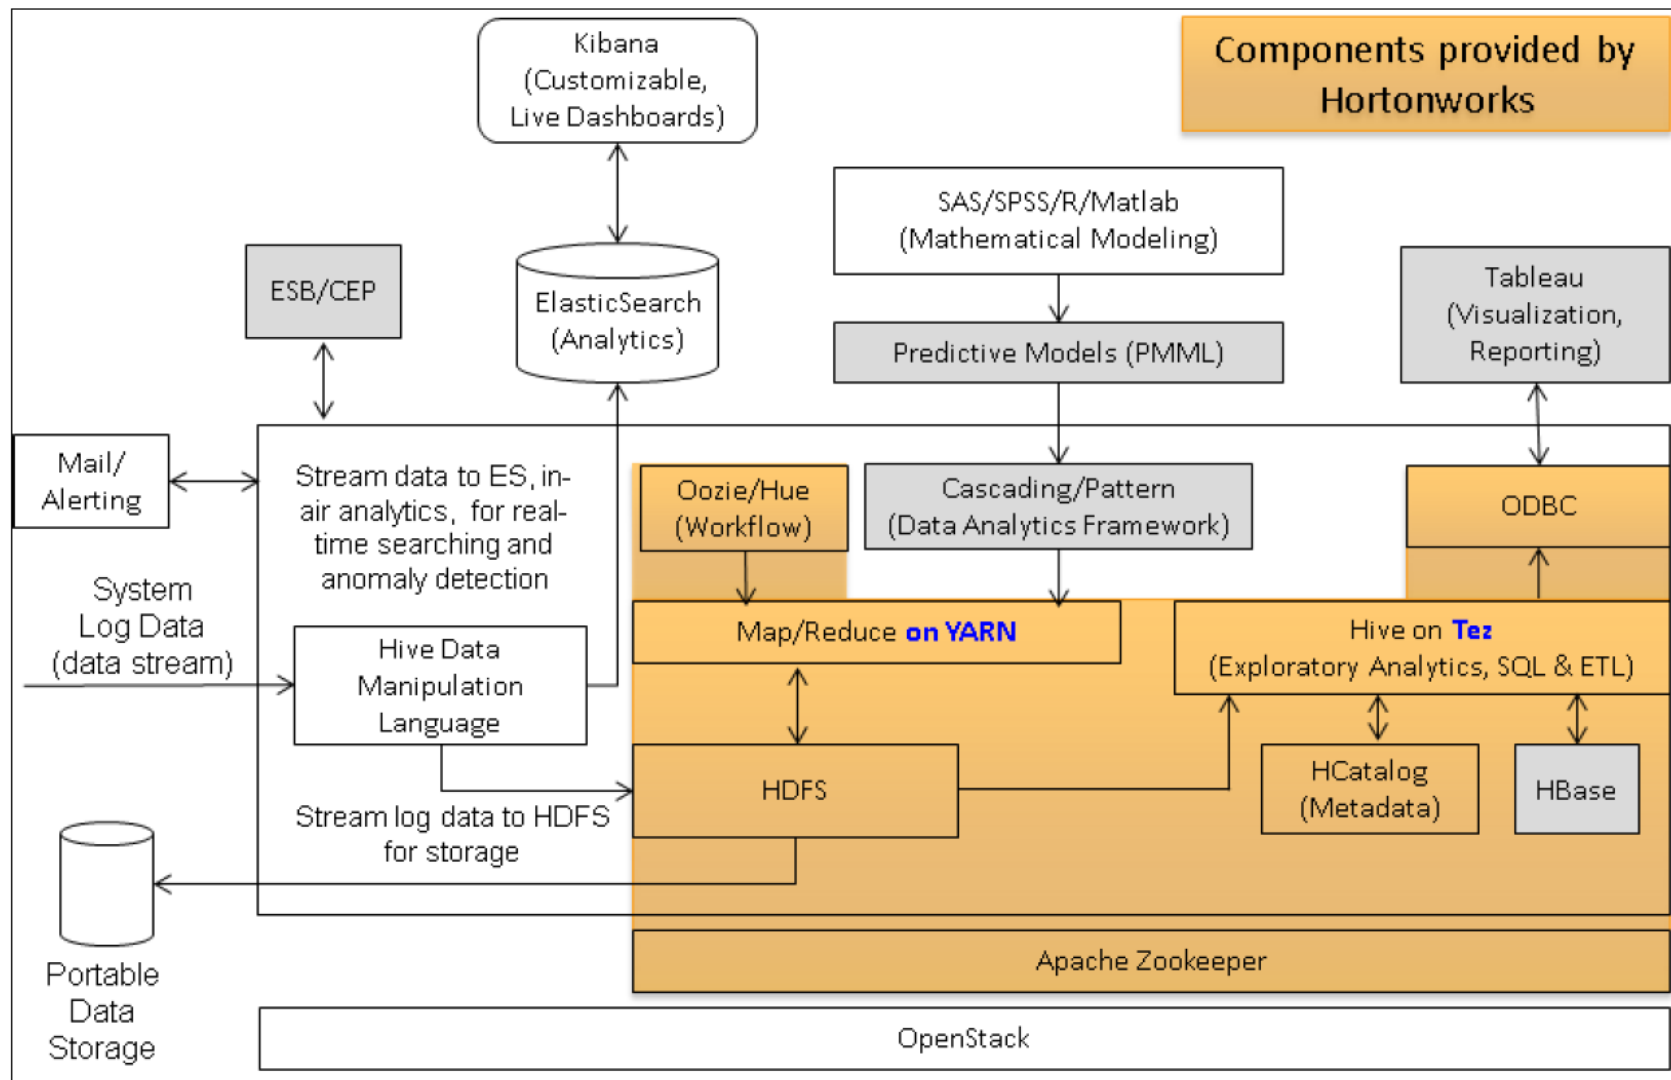
# Hadoop as a basic Big Data Platform



Source: Electronic Design, 2012, and Hortonworks, 2014

Source: pwc: <http://www.pwc.com/us/en/technologyforecast/2014/cloud-computing/assets/pdf/pwc-technologyforecast-data-lakes.pdf>


# Example Big Data Architecture




Boci, E. & Thistlethwaite, S.: A novel big data architecture in support of ADS-B data analytic  
*Proc. Integrated Communication, Navigation, and Surveillance Conference (ICNS), 2015, C1-1-C1-8*

# Architectures for Hadoop with SQL


## Hive (Native Hadoop)



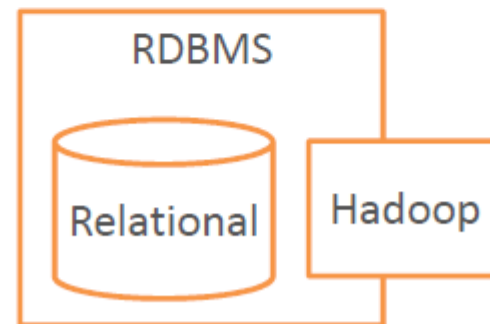
## Pure Hadoop SQL Engines



## Format-agnostic SQL Engines

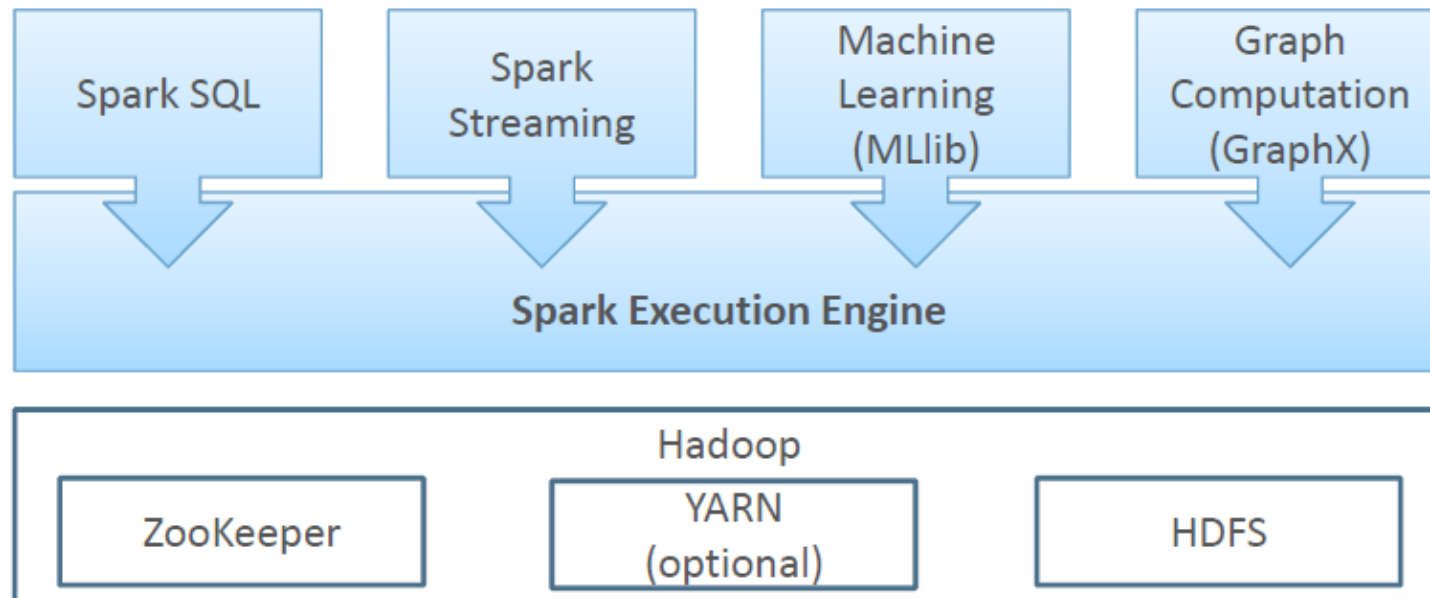


## RDBMS with Hadoop Access

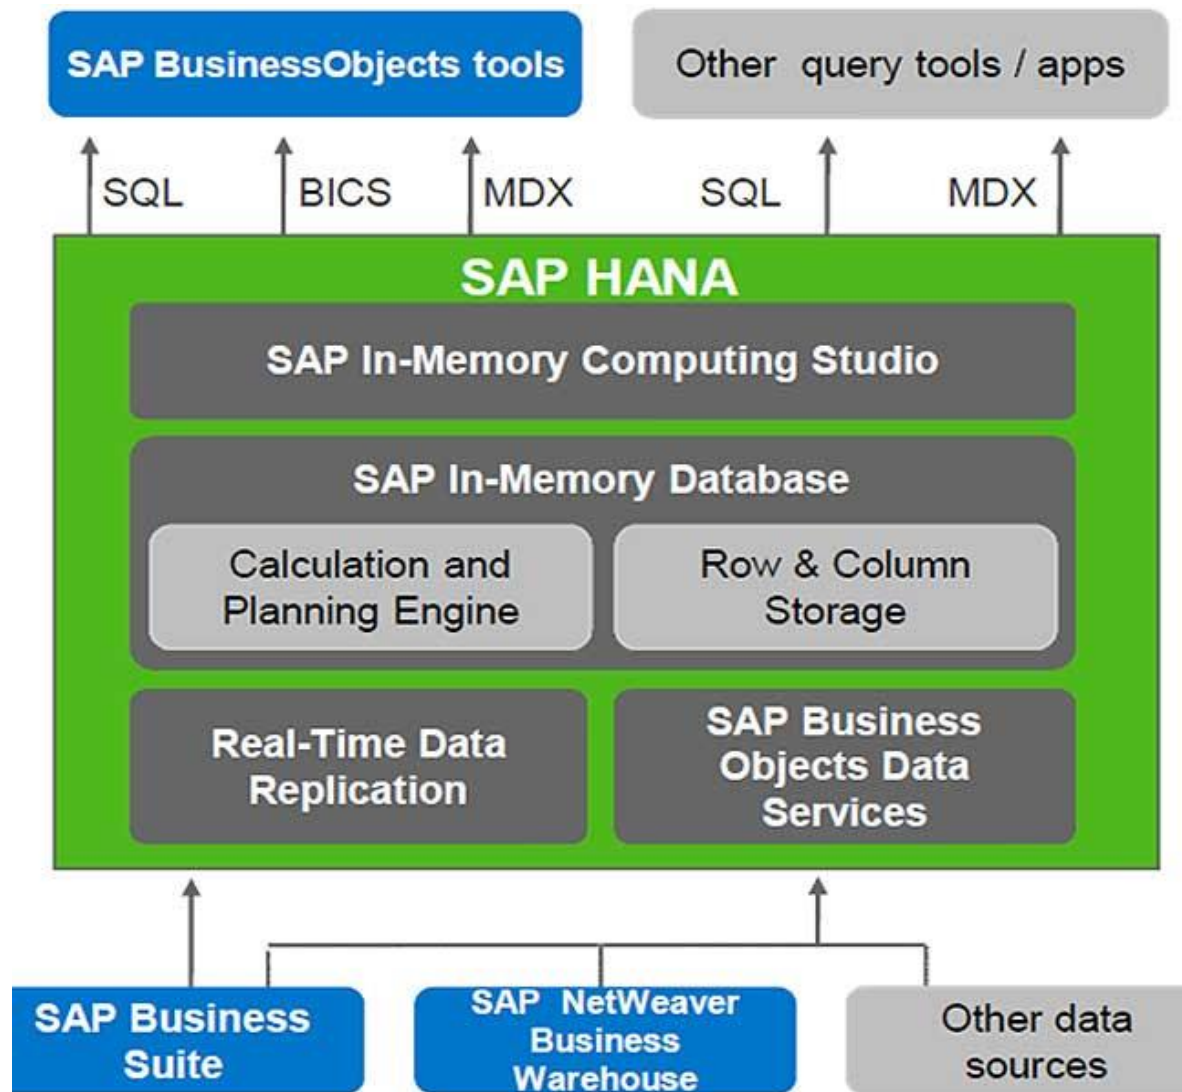


# In-Memory Database Systems

- Distributed data processing, but do as much as possible in-memory and avoid I/O operations (to disk or distributed file system)
- Example: Apache Spark
  - Distributed In-Memory Computing Framework
  - General framework for all kinds of SQL and non-SQL analytics



# Another Example for an In-Memory Database System




<https://www.stechies.com/overview-sap-hana-database-architecture/>



# Summary

- Main goal of DBMS: Data independence
- Gap in memory hierarchy is the challenge for a DBMS
- Five layered architecture
- Transaction management required for multi-user access
- Distributed Database Systems aim at higher reliability, availability, and scalability



# Review Questions

- Explain the concept of data independence!
- What are the layers of the five layered DBMS architecture?
- How is a query processed in a DBMS?
- Why is transaction management not assigned to a single layer in the DBMS architecture?
- What are the four levels in the IRDS architecture?
- What is distribution transparency in distributed database systems?
- What are basic characteristics of Big Data?

[Härder und Rahm, 2001] Härder, T. & Rahm, E. Datenbanksysteme: Konzepte und Techniken der Implementierung Springer Heidelberg, 2001, 2

[IRDS Framework Standard, ISO 10027:1990] <https://www.iso.org/obp/ui/#iso:std:iso-iec:10027:ed-1:v1:en>

[Özsu & Valduriez, 2011] Özsu, M. T. & Valduriez, P. Principles of distributed database systems Springer, 2011